

```
In [1]: import pandas as pd
```

```
In [2]: import numpy as np
```

```
In [3]: import matplotlib.pyplot as plt
```

```
In [4]: import seaborn as sns
```

```
In [5]: %matplotlib inline
```

```
In [6]: df = pd.read_csv("Housingdata.csv")
```

```
In [7]: df.head()
```

```
Out[7]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
0	0.00632	18.0	2.31	1	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	19.0	7.07	1	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	20.0	7.07	1	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	21.0	2.18	1	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	22.0	2.18	1	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	3.94	36.2

```
In [8]: df.tail()
```

```
Out[8]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
501	0.06263	289.0	11.93	496	0.573	6.593	69.1	2.4786	1	273	21.0	391.99	12.00	22.4
502	0.04527	290.0	11.93	497	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	9.08	20.6
503	0.06076	291.0	11.93	498	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	5.64	23.9
504	0.10959	292.0	11.93	499	0.573	6.794	92.0	2.3889	1	273	21.0	393.45	6.48	22.0
505	0.04741	293.0	11.93	500	0.573	6.030	93.0	2.5050	1	273	21.0	396.90	7.88	11.9

```
In [9]: print("The shape of the data is: ")
        df.shape
```

The shape of the data is:

```
Out[9]: (506, 14)
```


```
In [10]: df.isnull().sum()
```


```
Out[10]: CRIM      0
          ZN        0
          INDUS    0
          CHAS     0
          NOX      0
          RM       0
          AGE      0
          DIS      0
          RAD      0
          TAX      0
          PTRATIO  0
          B        0
          LSTAT    0
          MEDV     0
          dtype: int64
```

```
In [11]: corr = df.corr()
          corr.shape
```

```
Out[11]: (14, 14)
```

```
In [14]: plt.figure(figsize=(20, 20))  
sns.heatmap(corr, cbar=True, square=True, fmt='.1f', annot=True, annot_kws={'size': 15}, cmap='Blues')  
plt.show()
```





```
In [15]: x = df.drop(['MEDV'], axis = 1)
         y = df['MEDV']
```

```
In [16]: from sklearn.model_selection import train_test_split
         xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size = 0.2, random_state = 0)
```

```
In [17]: import sklearn
         from sklearn.linear_model import LinearRegression
         lm = LinearRegression()
```

```
In [18]: model = lm.fit(xtrain, ytrain)
```

```
In [19]: xtrain
```

Out[19]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
220	0.35809	126.0	6.20	215	0.507	6.951	88.5	2.8617	8	307	17.4	391.70	9.71
71	0.15876	16.5	10.81	66	0.413	5.961	17.5	5.2873	4	305	19.2	376.94	9.88
240	0.11329	30.0	4.93	235	0.428	6.897	54.3	6.3361	6	300	16.6	391.25	11.38
6	0.08829	12.5	7.87	1	0.524	6.012	66.6	5.5605	5	311	15.2	395.60	12.43
417	25.94060	205.0	18.10	412	0.679	5.304	89.1	1.6475	24	666	20.2	127.36	26.64
...
323	0.28392	111.0	7.38	318	0.493	5.708	74.3	4.7211	5	287	19.6	391.13	11.74
192	4.12579	98.0	3.44	187	0.437	7.178	26.3	6.4798	5	398	15.2	390.49	2.87
117	0.15098	23.0	10.01	112	0.547	6.021	82.6	2.7474	6	432	17.8	394.51	10.30
47	0.22927	82.0	17.96	42	0.448	6.030	85.5	5.6894	3	233	17.9	392.74	18.80
172	0.13914	78.0	4.05	167	0.510	5.572	88.5	2.5961	5	296	16.6	396.90	14.69

404 rows × 13 columns

```
In [20]: ytrain_pred=lm.predict(xtrain)
         ytest_pred=lm.predict(xtest)
```

```
In [21]: testdata=[[0.00632,18.0,2.31,0.0,0.538,6.575,65.2,4.0900,1.0,296.0,15.3,396.90,4.98]]
```

```
In [22]: test_pred = lm.predict(testdata)
         test_pred
```

```
/home/pratiksha/.local/lib/python3.8/site-packages/sklearn/base.py:465: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
  warnings.warn(
```

```
Out[22]: array([30.73465291])
```

```
In [23]: df1=pd.DataFrame(ytrain_pred,ytrain)
df2=pd.DataFrame(ytest_pred,ytest)
df1
```

```
Out[23]:
```

	0
--	---

MEDV	
------	--

26.7	31.083610
------	-----------

21.7	23.425128
------	-----------

22.0	28.744489
------	-----------

22.9	24.368941
------	-----------

10.4	6.805201
------	----------

...	...
-----	-----

18.5	19.764742
------	-----------

36.4	33.602929
------	-----------

19.2	24.152525
------	-----------

16.6	19.542249
------	-----------

23.1	23.475406
------	-----------

404 rows × 1 columns

```
In [24]: from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(ytest, ytest_pred)
print('MSE on test data:',mse)
```

```
mse1 = mean_squared_error(ytrain_pred, ytrain)
print('MSE on training data:',mse1)
```

MSE on test data: 36.514221822009354


MSE on training data: 20.6807311647163

```
In [25]: r2 = lm.score(xtest, ytest)
rmse = (np.sqrt(mean_squared_error(ytest, ytest_pred)))
print('r-squared: {}'.format(r2))
print('-----')
print('root mean squared error: {}'.format(rmse))
```

r-squared: 0.5515790030771384

root mean squared error: 6.042699878531893

```
In [26]: plt.scatter(ytrain ,ytrain_pred,c='blue',marker='o',label='Training data')
plt.scatter(ytest,ytest_pred ,c='lightgreen',marker='s',label='Test data')
plt.xlabel('True values')
plt.ylabel('Predicted')
plt.title("True value vs Predicted value")
plt.legend(loc= 'upper left')
plt.plot()
plt.show()
```

In []: